Trees & Forests

02/18/19

Andreas C. Müller

(Adapted and modified for CC 6021236 @ PCC/Ciencias/UCV by Eugenio Scalise, September 2019)

Why Trees?

- Very powerful modeling method – non-linear!
- Doesn’t care about scaling of distribution of data!
- “Interpretable”
- Basis of very powerful models!

Decision Trees for Classification

Idea: series of binary questions

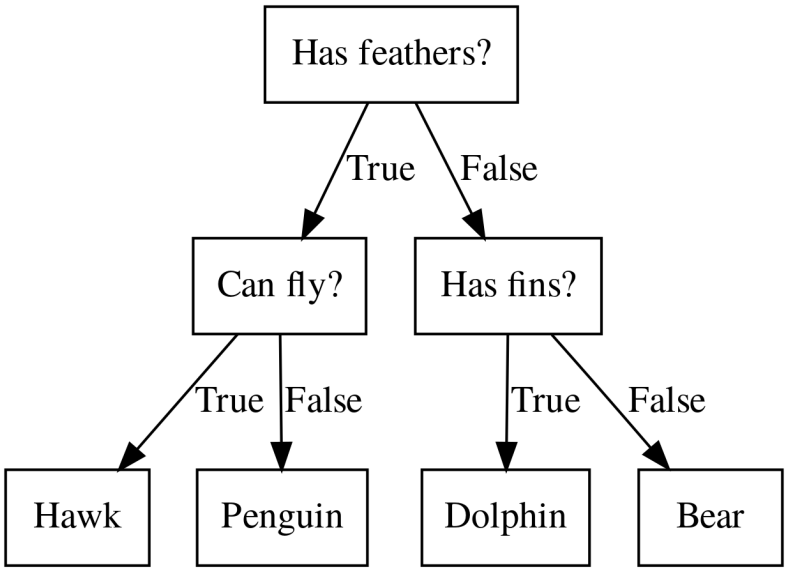

Building Trees

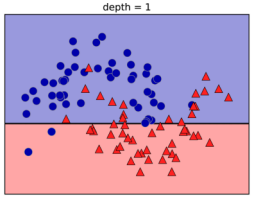

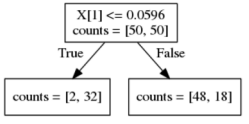

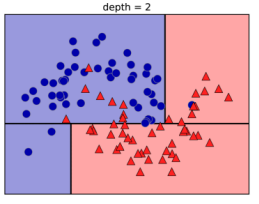

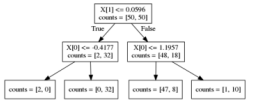

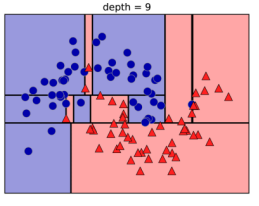

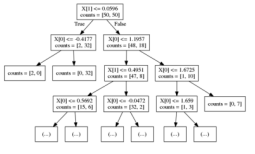

Continuous features:

- “questions” are thresholds on single features.
- Minimize impurity

Prediction

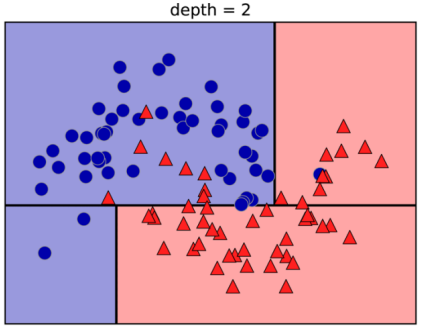
depth = 2

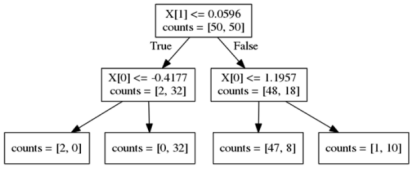
X[1] <= 0.0596
counts = [50, 50]
True
False
X[0] <= -0.4177
counts = [2, 32]
X[0] <= 1.1957
counts = [48, 18]
counts = [2, 0]
counts = [0, 32]
counts = [47, 8]
counts = [1, 10]

Visualizing trees with sklearn

Visualizing trees with sklearn

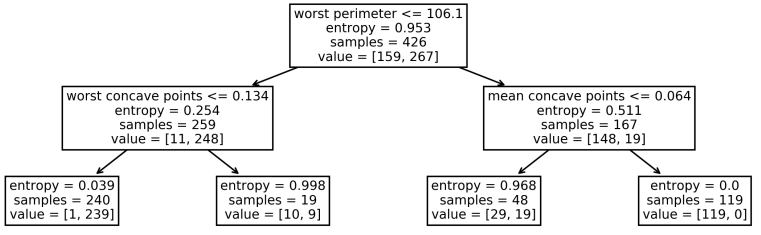

Parameter Tuning

- Limit tree size (pick one, maybe two):
  - max_depth
  - max_leaf_nodes
  - min_samples_split
  - min_impurity_decrease
  - ...

No pruning

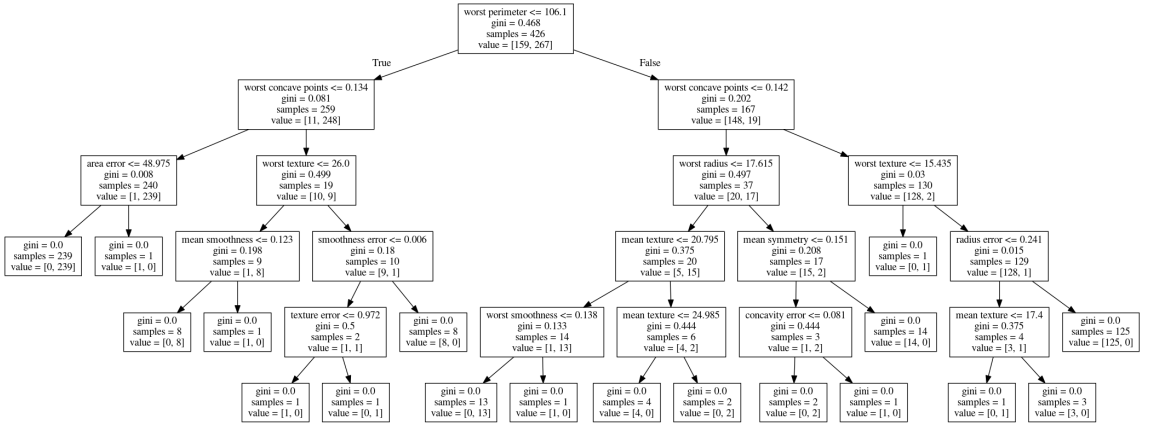
worst perimeter <= 106.1
gini = 0.468
samples = 426
value = [159, 267]
True
False
worst concave points <= 0.134
gini = 0.081
samples = 259
value = [11, 248]
worst concave points <= 0.142
gini = 0.202
samples = 167
value = [148, 19]
area error <= 48.975
gini = 0.008
samples = 240
value = [1, 239]
worst texture <= 26.0
gini = 0.499
samples = 19
value = [10, 9]
worst radius <= 17.615
gini = 0.497
samples = 37
value = [20, 17]
worst texture <= 15.435
gini = 0.03
samples = 130
value = [128, 2]
gini = 0.0
samples = 239
value = [0, 239]
gini = 0.0
samples = 1
value = [1, 0]
mean smoothness <= 0.123
gini = 0.198
samples = 9
value = [1, 8]
smoothness error <= 0.006
gini = 0.18
samples = 10
value = [9, 1]
mean texture <= 20.795
gini = 0.375
samples = 20
value = [5, 15]
mean symmetry <= 0.151
gini = 0.208
samples = 17
value = [15, 2]
gini = 0.0
samples = 1
value = [0, 1]
radius error <= 0.241
gini = 0.015
samples = 129
value = [128, 1]
gini = 0.0
samples = 8
value = [0, 8]
gini = 0.0
samples = 1
value = [1, 0]
texture error <= 0.972
gini = 0.5
samples = 2
value = [1, 1]
gini = 0.0
samples = 8
value = [8, 0]
worst smoothness <= 0.138
gini = 0.133
samples = 14
value = [1, 13]
mean texture <= 24.985
gini = 0.444
samples = 6
value = [4, 2]
concavity error <= 0.081
gini = 0.444
samples = 3
value = [1, 2]
gini = 0.0
samples = 14
value = [14, 0]
mean texture <= 17.4
gini = 0.375
samples = 4
value = [3, 1]
gini = 0.0
samples = 125
value = [125, 0]
gini = 0.0
samples = 1
value = [1, 0]
gini = 0.0
samples = 1
value = [0, 1]
gini = 0.0
samples = 13
value = [0, 13]
gini = 0.0
samples = 1
value = [1, 0]
gini = 0.0
samples = 4
value = [4, 0]
gini = 0.0
samples = 2
value = [0, 2]
gini = 0.0
samples = 2
value = [0, 2]
gini = 0.0
samples = 1
value = [1, 0]
gini = 0.0
samples = 1
value = [0, 1]
gini = 0.0
samples = 3
value = [3, 0]

max_depth = 4

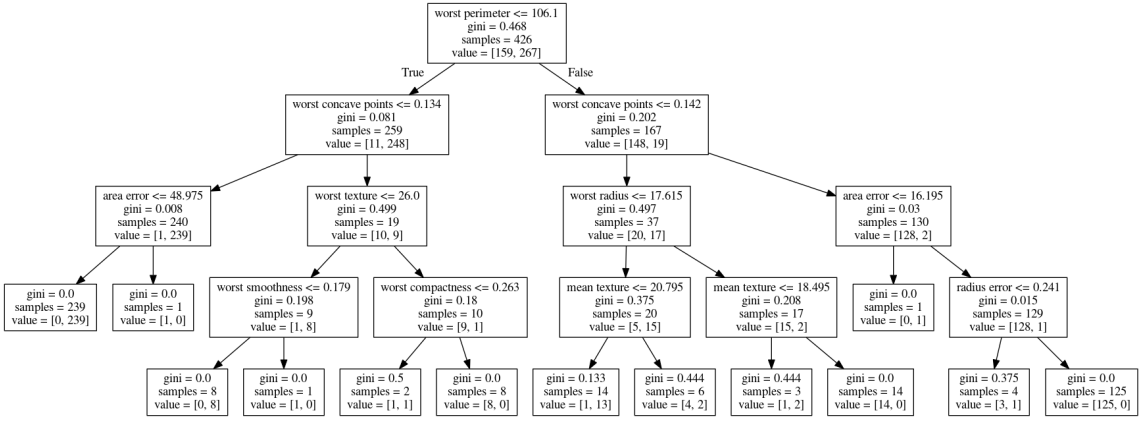
worst perimeter <= 106.1
gini = 0.468
samples = 426
value = [159, 267]
True
False
worst concave points <= 0.134
gini = 0.081
samples = 259
value = [11, 248]
worst concave points <= 0.142
gini = 0.202
samples = 167
value = [148, 19]
area error <= 48.975
gini = 0.008
samples = 240
value = [1, 239]
worst texture <= 26.0
gini = 0.499
samples = 19
value = [10, 9]
worst radius <= 17.615
gini = 0.497
samples = 37
value = [20, 17]
area error <= 16.195
gini = 0.03
samples = 130
value = [128, 2]
gini = 0.0
samples = 239
value = [0, 239]
gini = 0.0
samples = 1
value = [1, 0]
worst smoothness <= 0.179
gini = 0.198
samples = 9
value = [1, 8]
worst compactness <= 0.263
gini = 0.18
samples = 10
value = [9, 1]
mean texture <= 20.795
gini = 0.375
samples = 20
value = [5, 15]
mean texture <= 18.495
gini = 0.208
samples = 17
value = [15, 2]
gini = 0.0
samples = 1
value = [0, 1]
radius error <= 0.241
gini = 0.015
samples = 129
value = [128, 1]
gini = 0.0
samples = 8
value = [0, 8]
gini = 0.0
samples = 1
value = [1, 0]
gini = 0.5
samples = 2
value = [1, 1]
gini = 0.0
samples = 8
value = [8, 0]
gini = 0.133
samples = 14
value = [1, 13]
gini = 0.444
samples = 6
value = [4, 2]
gini = 0.444
samples = 3
value = [1, 2]
gini = 0.0
samples = 14
value = [14, 0]
gini = 0.375
samples = 4
value = [3, 1]
gini = 0.0
samples = 125
value = [125, 0]

max_leaf_nodes = 8

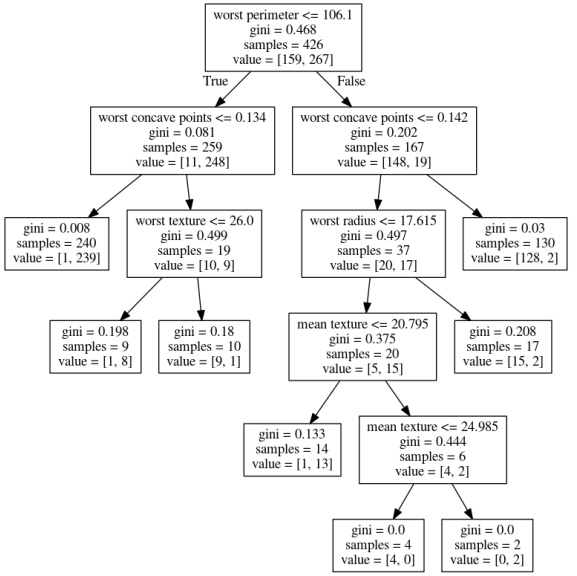

min_samples_split = 50

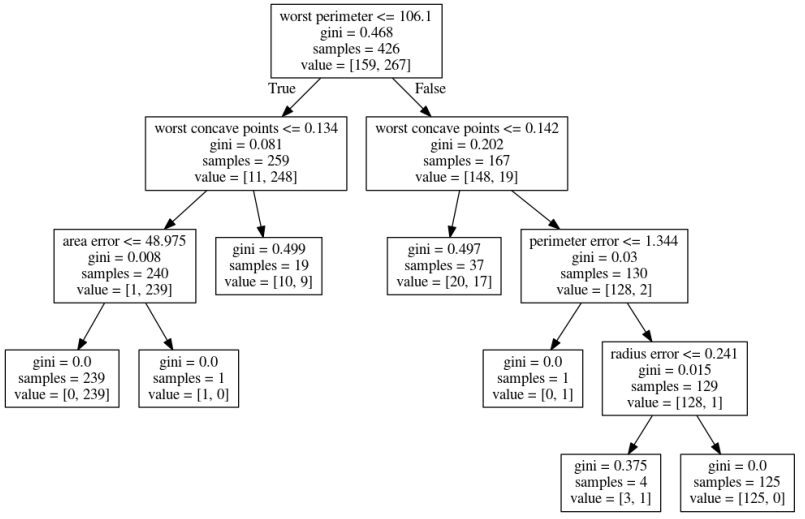

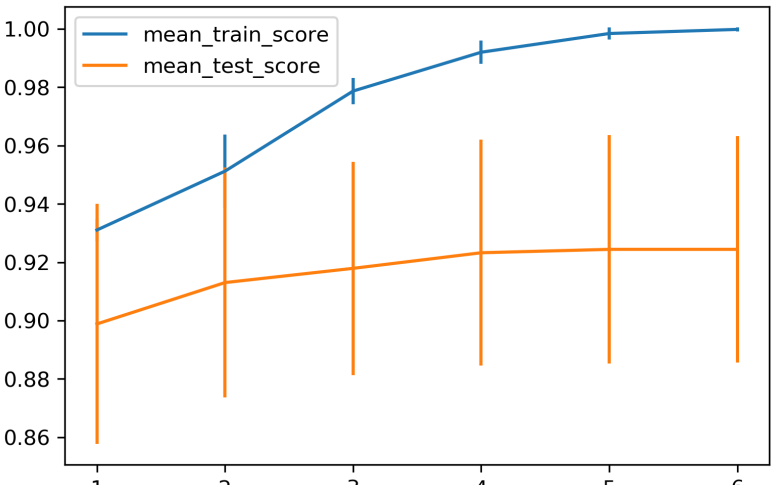
mean_train_score
mean_test_score
1.00
0.98
0.96
0.94
0.92
0.90
0.88
0.86

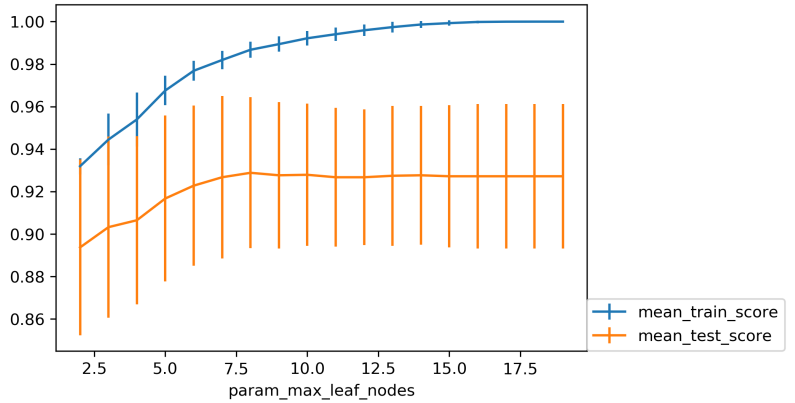
1.00
0.98
0.96
0.94
0.92
0.90
0.88
0.86
2.5
5.0
7.5
10.0
12.5
15.0
17.5
param_max_leaf_nodes
mean_train_score
mean_test_score

Instability

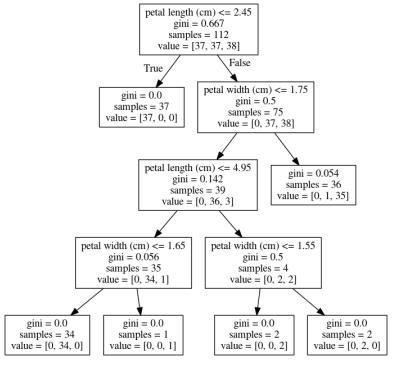

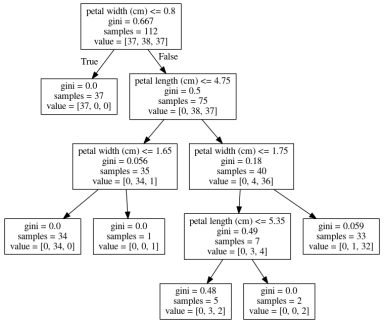

Feature importance

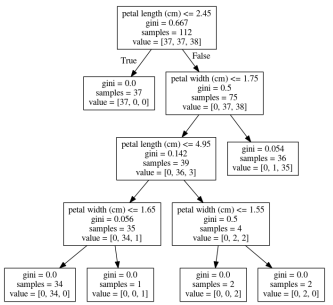

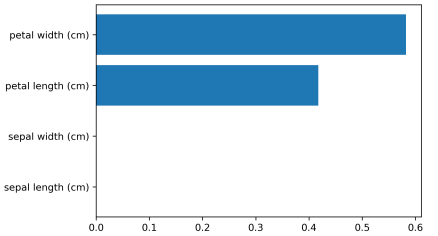

Ensemble Models (Random Forests)

Poor man's ensembles

- Build different models
- Average the result
- More models are better – if they are not correlated.
- Also works with neural networks
- You can average any models as long as they provide calibrated ("good") probabilities.
- Scikit-learn: VotingClassifier

Bagging (Bootstrap AGGregation)

- Generic way to build “slightly different” models
- BaggingClassifier, BaggingRegressor

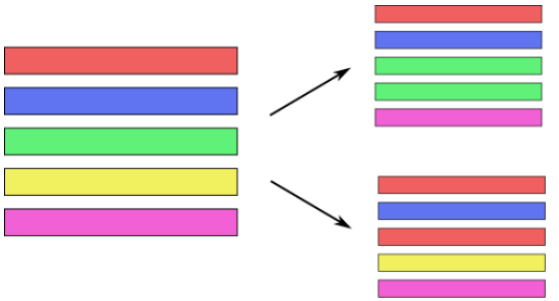

Random Forests

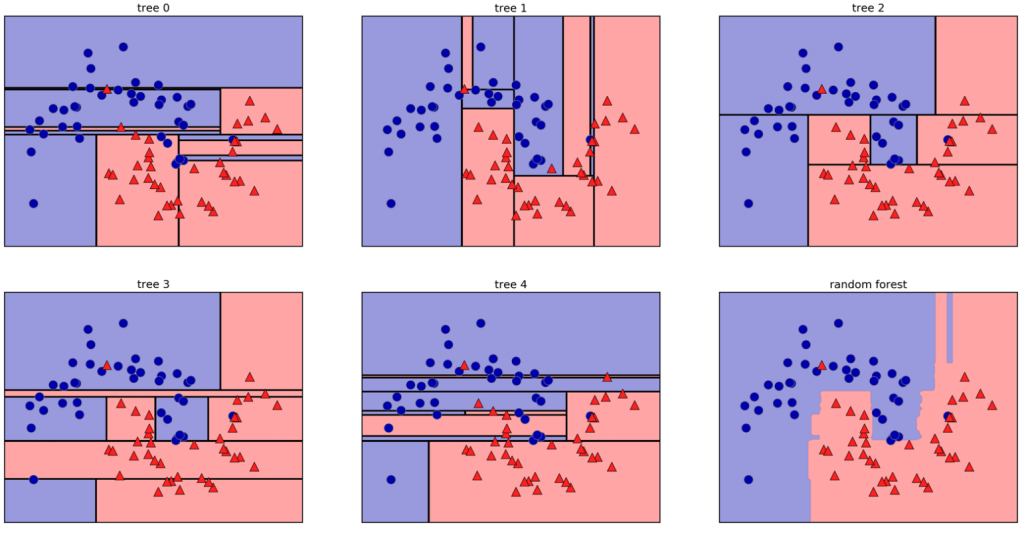

Randomize in two ways

- For each tree:
  - Pick bootstrap sample of data
- For each split:
  - Pick random sample of features
- More trees are always better

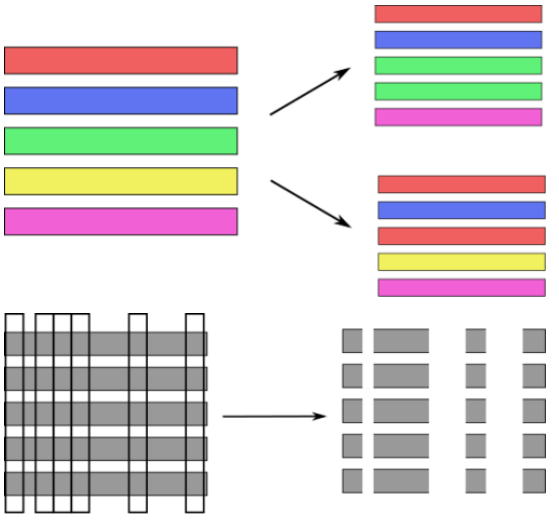

Tuning Random Forests

- Main parameter: max_features
  - around sqrt(n_features) for classification
  - Around n_features for regression
- n_estimators > 100
- max_depth, max_leaf_nodes, min_samples_split again